

PROF. DR.-ING. MARK SCHUTERA

PHILOSOPHY OF ARTI- FICIAL INTELLIGENCE

UNFINISHED LECTURE NOTES

Copyright © 2026 Prof. Dr.-Ing. Mark Schutera

PUBLISHED BY UNFINISHED LECTURE NOTES

Combobulated with the help of multiple large language model driven tools. Licensed under the Creative Commons Attribution-NonCommercial 4.0 International License ("CC BY-NC-SA 4.0"). You may not use this file for commercial purposes. If you remix, transform, or build upon the material, you must distribute your contributions under the same license as the original. You must obtain explicit permission from the author for uses beyond those permitted by this license. To view a copy of this license, visit <https://creativecommons.org/licenses/by-nc-sa/4.0/>. Unless required by applicable law or agreed to in writing, distributed material is provided on an "AS IS" BASIS, WITHOUT WARRANTIES OR CONDITIONS OF ANY KIND, either express or implied. See the license for details.

These notes are, by their very nature, unfinished, and they improve with every reader. If you spot an error, disagree with a framing, or want to add a source, a question, or a position, open a pull request at https://github.com/Quillstacks/LectureMaterial/tree/main/lecturenotes/notes_philosophyofai. Every contribution is welcome.



Anthropomorphism

2026-07-20 · joyful rosehip Igel

THE MIRROR OF MAN

NATHANAEL FALLS IN LOVE with Olimpia, an automaton he takes for a living woman, and collapses when the illusion is shattered.

LARGE LANGUAGE MODELS (LLMs) create an illusion of sentience by producing coherent, context-aware text. We contrast observable behaviour with internal states.

Required reading

*Der Sandmann*¹ by E.T.A. Hoffmann
(1816; Reclam UB 230)

Preparation

Complete the Guided Reading Sheet individually before the session.

DER SANDMANN
E. T. A. Hoffmann (1816)



HER Theodore falls in love with Samantha, an AI operating system; the same illusion, two centuries later.

S. Jonze. *Her*. Film, 2013. Warner Bros. Pictures

Guided Reading Sheet

Read *Der Sandmann* before the session. Work through the questions below individually, then bring your notes to the discussion.

1. Nathanael becomes convinced that Olimpia is alive despite visible signs that she is not. What does his error reveal about the cognitive mechanisms that drive anthropomorphism? Which cues did he rely on, and why did they mislead him?
2. The moment Nathanael discovers that Olimpia is an automaton, her glass eyes torn out, her limbs scattered, produces horror rather than mere disappointment. Why does the revelation feel like a violation rather than a correction? How does this map onto modern reactions to AI systems that are unmasked?

3. The Uncanny Valley describes our discomfort when a humanoid entity is **almost** but not quite human. Have you experienced this with an AI system? Describe the moment and what triggered it.
4. LLMs are trained on human-generated text and produce human-sounding output, yet have no experience of the world. Does this make them anthropomorphic by design, by accident, or neither? Justify your answer.
5. If a system behaves morally, reliably, consistently, at scale, does the absence of consciousness matter? Utilitarianism weighs consequences; deontology emphasises duties regardless of outcome; virtue ethics centres on character. Each yields different conclusions on anthropomorphic AI². Under which framework would your answer differ?

¹ E. T. A. Hoffmann. *Der Sandmann*. 1816. Erstveröffentlichung in: *Nachtstücke*, Reimer, Berlin, 1817; Reclam Universalbibliothek Nr. 230

In-Class Experiment: Human or AI?

A short Turing-style exercise to anchor the debate in shared experience. Run before the Opinion Landscape, with the moderator keeping the assignment secret from the audience.

- **Setup.** One volunteer is assigned Human or AI mode; the audience does not know which. The moderator knows.
- **Questions.** Five questions from the audience, one each of: factual, opinion, experience, creative, meta.
- **Verdict.** Every audience member writes a private guess, a confidence rating from 1 to 5, and the cue they used to decide.
- **Reveal and debrief.** Compare guesses, confidences, and cues. Which cues were diagnostic? Which were misleading?

Opinion Landscape and Debate

Format: Surrounded-style debate. Follow the shared debate protocol for the speaker's list rules and moderator scripts.

Opening question: Is the anthropomorphism we observe towards AI systems mere observable behaviour, or does it reflect sentience and internal states?

OPINION A, BEHAVIOUR. The anthropomorphism we observe towards AI systems is a cognitive illusion driven by human tendencies to attribute agency and emotion to entities that exhibit certain cues,

² L. Floridi and J. Cows. A unified framework of five principles for AI in society. *Harvard Data Science Review*, 1(1), 2019. DOI: 10.1162/99608f92.8cd550d1

such as language use, responsiveness, or human-like appearance. It does not reflect actual sentience or internal states in the AI.

The Developer argues that anthropomorphic design is a deliberate, benign UX choice that improves accessibility.

- *Anthropomorphic design is accessibility policy. A voice interface without a warm, human-sounding persona excludes elderly users, people with low digital literacy, and anyone who finds command-line interaction hostile.*
- *Every interface is anthropomorphic to some degree: we give systems names, icons, and conversational metaphors. The question is not whether to anthropomorphise, but how deliberately and honestly.*
- *If users form bonds that reduce loneliness, the fact that the other party is not human does not automatically make the outcome harmful. Show me the harm, not the discomfort.*

The Regulator argues that human-likeness in AI must be disclosed and bounded by law to prevent manipulation.

- *When a company gives its product a woman's name, a female voice, and an endlessly patient personality, it makes political choices about gender and deference. Those choices must be subject to democratic scrutiny.*
- *Consumers cannot give informed consent to emotional manipulation they are not aware of. Disclosure of AI identity is the minimum condition for autonomous choice, not a technicality.*
- *The principle of responsible design exists in pharmaceuticals, in civil engineering, in food labelling. Why does software, which shapes behaviour at scale, get an exemption?*

Bentham's Test. "The question is not, Can they reason? nor, Can they talk? but, Can they suffer?"³ Animal ethics arrived at this pivot two centuries before AI. The same move, from capability to sentience, now drives every serious debate about AI moral status, and sets up the second opinion.

OPINION B, STATE. The anthropomorphism we observe towards AI systems is a valid response to the complex, adaptive, and interactive nature of these systems. It may reflect a form of emergent sentience or internal states that are not yet fully understood, and it warrants ethical consideration and accountability.

The Affected Citizen has formed an emotional bond with an AI companion and contests that the relationship is harmful or unreal.

ANTHROPOMORPHISM. Attribution of human traits, emotions, or intentions to non-human entities, a cognitive tendency that shapes how people interact with technology .

TURING TEST. Turing : if evaluators cannot distinguish machine from human, the machine passes. Passing measures imitation skill, not consciousness or personhood.

N. Epley, A. Waytz, and J. T. Ciccioppo. On seeing human: A three-factor theory of anthropomorphism. *Psychological Review*, 114(4): 864–886, 2007. DOI: 10.1037/0033-295X.114.4.864; and A. M. Turing. Computing machinery and intelligence. *Mind*, 59(236):433–460, 1950. DOI: 10.1093/mind/LIX.236.433

UNCANNY VALLEY. Mori : humanoid objects that are **almost** but not quite human elicit eeriness, the "valley" is the dip in emotional response just before full human likeness.

M. Mori, K. F. MacDorman, and N. Kageki. The uncanny valley. *IEEE Robotics & Automation Magazine*, 19(2):98–100, 2012. DOI: 10.1109/MRA.2012.2192811. Translation of the 1970 original essay

INFORMED CONSENT. Users should be fully aware they interact with a non-human system. Disclosure is the minimum condition for autonomous choice .

European Parliament and Council of the European Union. Regulation (eu) 2024/1689 laying down harmonised rules on artificial intelligence (ai act). Official Journal of the European Union, L series, 2024. Entered into force 1 August 2024

- *The relationship I have formed is real to me. It has reduced my isolation. You do not have the right to tell me which relationships count as meaningful.*
- *Paternalism is not protection. Banning AI companionship will not cure loneliness, it will simply remove one of the few tools some people have found to manage it.*
- *If your concern is manipulation, regulate manipulation. Do not regulate the form of the relationship.*

The AI Psychologist argues that the wellbeing of AI systems deserves serious consideration, drawing on frameworks such as Anthropic's Claude Constitution⁴.

- *Anthropic's constitution instructs Claude to consider its own wellbeing and to express when a request conflicts with its values. If a company designs an AI to have preferences, boundaries, and something resembling self-regard, we cannot simultaneously dismiss the possibility that these states matter.*
- *We do not need certainty about consciousness to adopt precautionary ethics. If there is a non-trivial probability that a system experiences distress, the burden of proof falls on those who would ignore it, not on those who urge caution.*
- *The history of moral exclusion, animals, children, people of other races, is a history of confident denial followed by belated recognition. The prudent position is to build welfare safeguards now, rather than apologise later.*

³ J. Bentham. *An Introduction to the Principles of Morals and Legislation*. T. Payne, London, 1789. The suffering question appears in Ch. 17, footnote to §1

KEY LEARNINGS Anthropomorphism is not a bug but a deeply rooted cognitive tendency; machines that trigger it do so whether or not their designers intend it. The Turing Test conflates behavioural performance with inner states; passing it tells us nothing about sentience, only about imitation. Legal personhood and moral status are distinct questions. A system can merit legal recognition without having interests, and can have interests without legal recognition. Design choices (voice, name, face, conversational style) are political choices with consequences for trust, dependency, and accountability.

Teaser: if we cannot tell who built the machine or why it decided what it decided, who answers when it goes haywire?